Vectorization

Oliver Heidmann

Arbeitsbereich Wissenschaftliches Rechnen
Fachbereich Informatik
Fakultät für Mathematik, Informatik und Naturwissenschaften
Universität Hamburg

2016-12-22

Gliederung (Agenda)

The Program:
Simulation/Game/Analytics which processes huge amounts of data.
It is already written in an data oriented style.

The Problem:
The execution time is way too high.

What can we do?

Steps of making code faster:

- manual optimizations
- parallelization
- buying better hardware
- buying more hardware

Steps of making code faster:

- manual optimizations
- parallelization
- => vectorization <=
- buying better hardware
- buying more hardware

Vectorization

What is Vectorization?

Vectorization allows us to compute multiple operations at once.

How is that possible?

- extended set of CPU instructions

Important differences between instruction sets:

- SSE(Streaming SIMD Extensions)
  - only single precision floats
  - 8 128-bit vector registers
  - first supported by intel pentium 3

Important differences between instruction sets:

- SSE(Streaming SIMD Extensions)
  - only single precision floats
  - 8 128-bit vector registers
  - first supported int intel pentium 3
- SSE2
  - added support for 16-bit short, 32-int, 64-double-precision and 64-int
  - added 8 new vector registers for x64

Important differences between instruction sets:

- SSE(Streaming SIMD Extensions)
  - only single precision floats
  - 8 128-bit vector registers
  - first supported int intel pentium 3
- SSE2
  - added support for 16-bit short, 32-int, 64-double-precision and 64-int
  - added 8 new vector registers for x64
- AVX/AVX2(Advanced Vector Extensions)
  - now 256-bit registers
  - added three-operand SIMDs
  - added gather support

Vectorization

What is Vectorization?

Vectorization allows us to compute multiple operations at once.

How is that possible?

- extended set of CPU instructions
- vector units

What are those units?

- special computation units
- every modern CPU implements them
- calculate multiple results from multiple inputs in one instruction

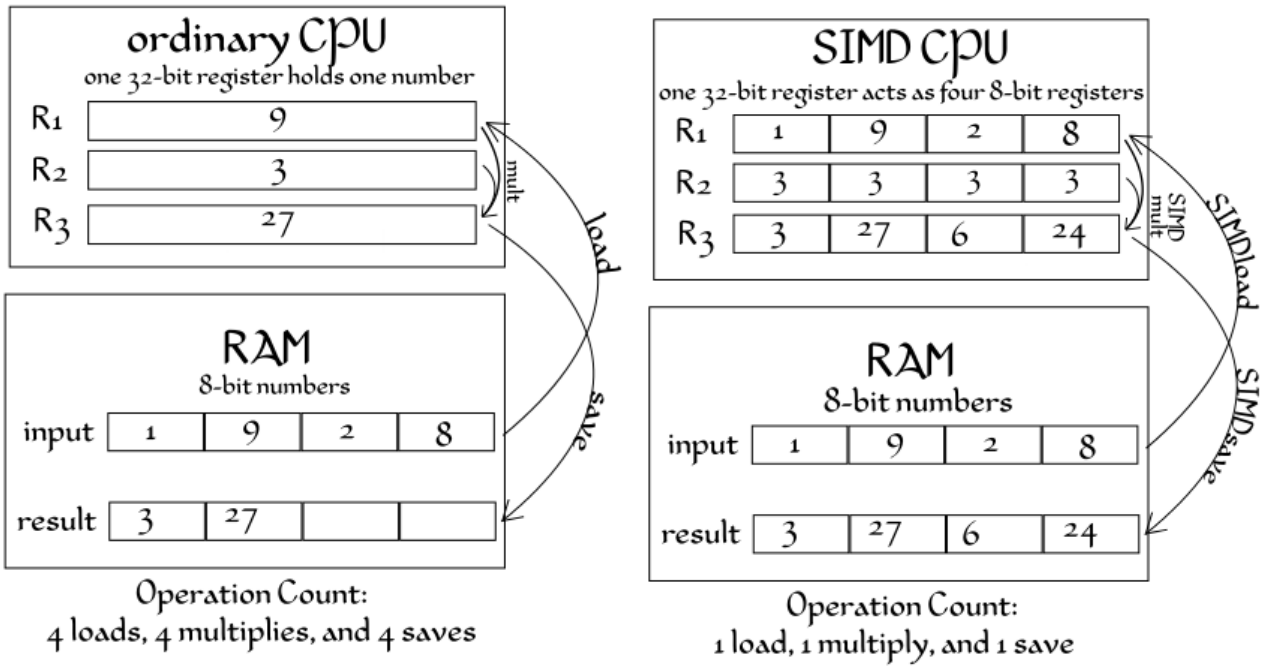

Vectorization

What is Vectorization?

Vectorization allows us to compute multiple operations at once.

How is that possible?

- extended set of CPU instructions
- vector units
- vector registers

Vector Registers

- extra registers on the CPU
- can store and load multiple values at once

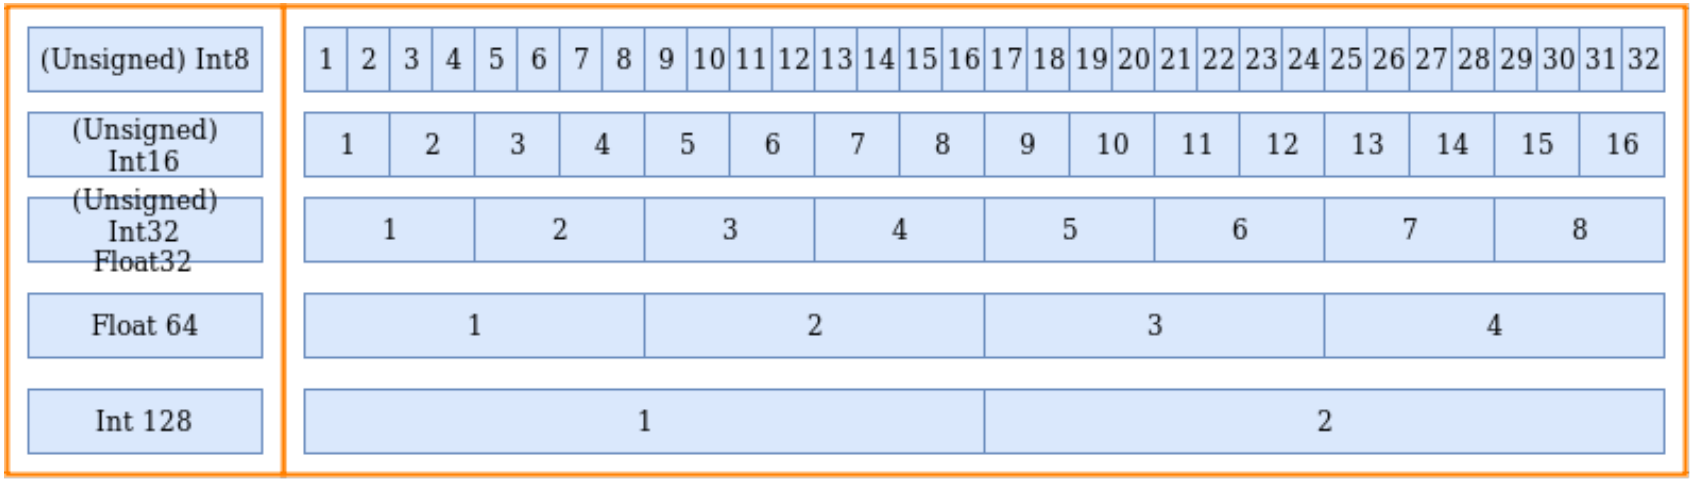

Vector Registers

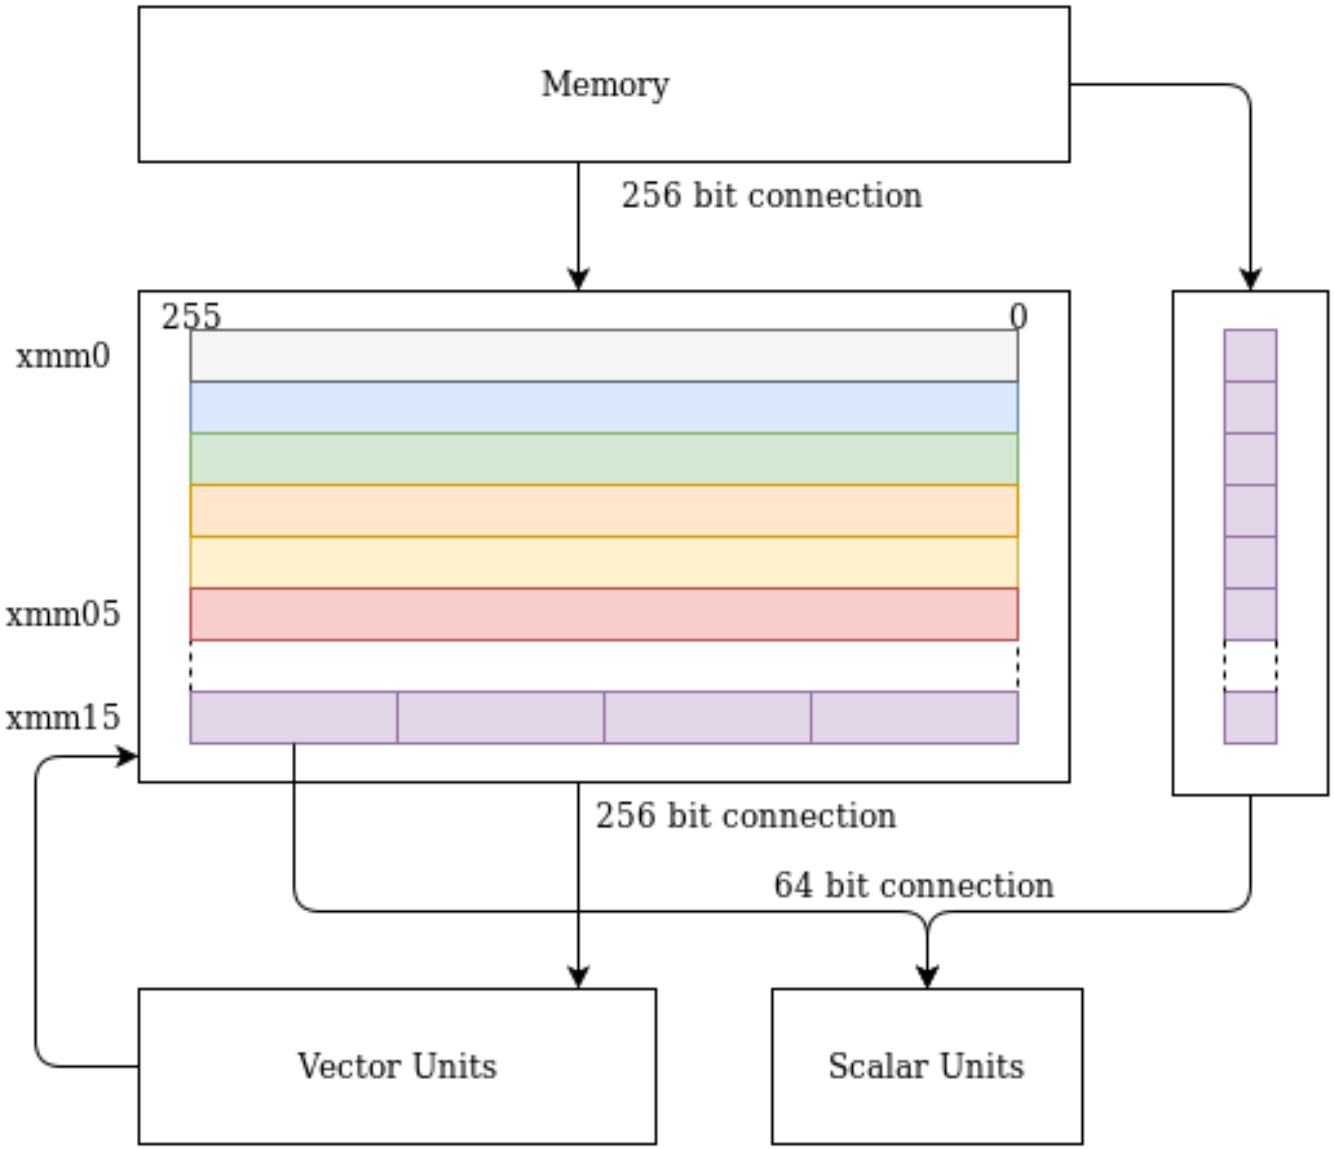

Vectorization

What is Vectorization?

Vectorization allows us to compute multiple operations at once.

How is that possible?

- extended set of CPU instructions
- vector units
- vector registers

Vectorization

What is Vectorization?

Vectorization allows us to compute multiple operations at once.

How is that possible?

- extended set of CPU instructions
- vector units
- vector registers
- everything implemented in silicon

Vectorization

What is Vectorization?

Vectorization allows us to compute multiple operations at once.

How is that possible?

- vector units
- vector registers
- extended set of CPU instructions
- everything implemented in silicon

What speedup can we expect?

type-width	128-bit	256-bit
8	1600%	3200%
16	800%	1600%
32	400%	800%
64	200%	400%

Real speedup will not be as huge

- overhead from loops
- cache misses/ memory access times
- data layout not perfect

Vectorization

What is Vectorization?

Vectorization allows us to compute multiple operations at once.

How is that possible?

- vector units
- vector registers
- extended set of CPU instructions
- everything implemented in silicon

The effect:

- huge speedups

How can I use vectorization?

The compiler does that for us if we tell him to.
Example for gcc:

- gcc standard optimizations do not vectorize
- -O3 enables auto vectorization
- -O3 does it by using the -ftree-vectorize flag
- -fopt-info-vec enables vectorization report
- -save-temps saves the temporary files eg. assembler code

What makes my code eligible for vectorization?

- calculations over arrays
- code must be in the innermost loop
- no if statements
- only inlined functions
- continuous data chunks

data organisation

Context: distance calculations; sqrt(x * x + y * y + z * z)

```
struct vector
{
    float x;
    float y;
    float z;
}
```

```
struct particle
{
    vector pos;
    vector velo;
    vector accel;
}
```

This wont work well

- data is not coherent

data organisation

Context: distance calculations; sqrt(x * x + y * y + z * z)

```
struct vectors
{
    float x[particle_cnt];
    float y[particle_cnt];
    float z[particle_cnt];
}
```

```
struct particles
{
  vectors pos;
  vectors velo;
  vectors accel;
}
```

This will work well

- data is now coherent

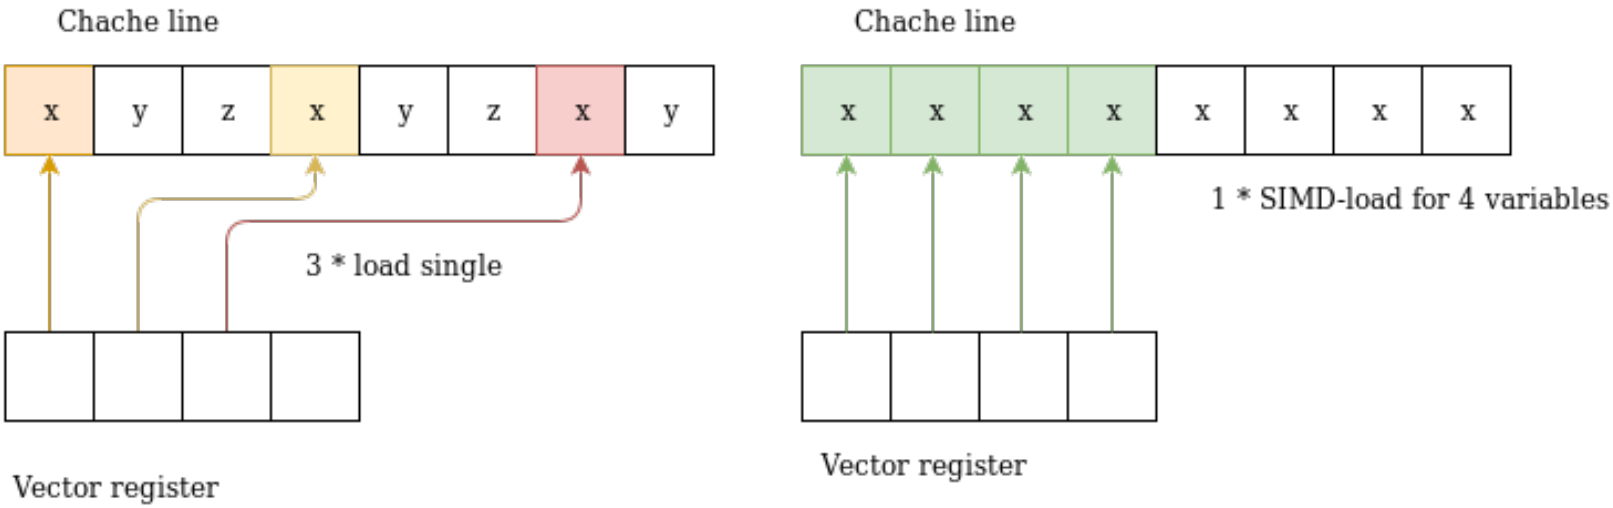
Chache line
x y z x y z x y
3 * load single
Vector register
Chache line
x x x x x x x x
1 * SIMD-load for 4 variables
Vector register

```
void test(float * vec1, float * vec2, float * res) {
    for (unsigned long i = 0; i < vector_size; i++) {
        res[i] += vec2[i] * vec1[i];
    }
}
```

- program checks for overlapping arrays parts
- program needs to check for aliasing

The restrict keyword:

Tells the compiler that the pointers are not aliased.
Meaning that the (sub)arrays are not overlapping or the same.

```
void test(float *__restrict vec1,
          float *__restrict vec2,
          float *__restrict res) {
    for (unsigned long i = 0; i < vector_size; i++) {
        res[i] += vec2[i] * vec1[i];
    }
}
```

- needs information about type boundaries

__attribute__((__aligned__(type_size))) :
Tells the compiler the size of the type in bit.
So that it is known how big a to be loaded bit word is.
Otherwise size will be checked at runtime.

```
constexpr size_t float_size = sizeof(float) * 8;
typedef float float_32 __attribute__((__aligned__(float_size)))

void test(float_32 *__restrict vec1,
          float_32 *__restrict vec2,
          float_32 *__restrict res)
{
    for (unsigned long i = 0; i < vector_size; i++) {
        res[i] += vec2[i] * vec1[i];
    }
}
```

Vectorization: Pros/Cons

Pros:

- depending on numeric type we can gain huge to immense speedup
- most modern systems support vectorization
- no extra cost for new hardware
- no extra software needed

Cons:

- complicated to implement for object oriented design
- exact result only visible in assembler code

Conclusion

- vectorization is a from of optimization
  - supported by modern compilers (gcc 4.6 and onward)
  - supported in modern hardware
  - when done right gives immense speedup
- Vectorizing
  - compiler does it for us
  - if it gets enough info
  - needs coherent data layout

- talk about vectorization by Ulrich Drepper
  `https://www.youtube.com/watch?v=DXPfE2jGqg0`
- talk about vectorization by James Reinders
  `https://www.youtube.com/watch?v=hyZMssi_gZY&t=1640s`
- Article about auto vectorization (caution! for gcc 4.7)
  `https://locklessinc.com/articles/vectorize/`

Literatur