

# Titel des Vortrags

Name des Autors

Arbeitsbereich Wissenschaftliches Rechnen  
Fachbereich Informatik  
Fakultät für Mathematik, Informatik und Naturwissenschaften  
Universität Hamburg

2015-01-01



Universität Hamburg  
DER FORSCHUNG | DER LEHRE | DER BILDUNG

**informatik**  
**die zukunft**

# Gliederung (Agenda)

- 1 The problem at hand
- 2 What is vectorization?
- 3 Vectorizing code
- 4 Conclusion
- 5 Literatur

The Program:

Simulation/Game/Analytics which processes huge amounts of data.  
It is already written in an data oriented style.

The Problem:

The execution time is way to high.

What can we do?

## Steps of making code faster:

- reduce cash misses
- manual optimizations
- parallelization
- reduce overhead
- buying better hardware
- buying more hardware

## Steps of making code faster:

- reduce cash misses
- manual optimizations
- parallelization
- $\Rightarrow$  vectorization  $\Leftarrow$
- reduce overhead
- buying better hardware
- buying more hardware

What does vectorization mean?

# Vectorization

What is Vectorization?

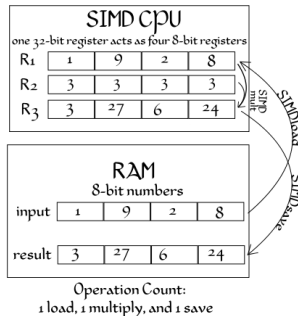
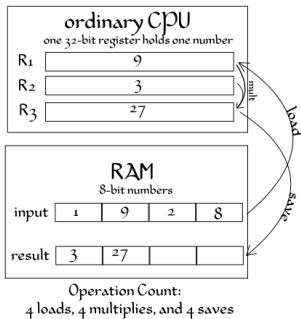
The use of a cpu's:

- vector units

What are those units?

# What are those units?

- special computation units
- every modern cpu implements them
- calculate multiple results from multiple inputs in one cycle





What are those units?

# Vectorization

What is Vectorization?

The use of a cpu's:

- vector units
- full vector registers

What are those units?

# Vector Registers

| Lanes per type in a 128-bit SIMD register |       |   |   |   |       |   |   |   |       |    |    |    |       |    |    |    |
|---|-------|---|---|---|-------|---|---|---|-------|----|----|----|-------|----|----|----|
| (U)Int8x16                                | 1     | 2 | 3 | 4 | 5     | 6 | 7 | 8 | 9     | 10 | 11 | 12 | 13    | 14 | 15 | 16 |
| (U)Int16x8                                | 1     |   | 2 |   | 3     |   | 4 |   | 5     |    | 6  |    | 7     |    | 8  |    |
| (U)Int32x4<br>Float32x4                   | 1 (x) |   |   |   | 2 (y) |   |   |   | 3 (z) |    |    |    | 4 (w) |    |    |    |
| Float64x2                                 | 1 (x) |   |   |   |       |   |   |   | 2 (y) |    |    |    |       |    |    |    |

What are those units?

# Vectorization

What is Vectorization?

The use of a cpu's:

- vector units
- full vector registers
- extended set of cpu instructions

# Extended vector instructions

MOVAPS(x,y)z

moves a memory line starting with x and with y size  
to vector register z

MOV(x,y)

moves single value from memory x into scalar register y

# Instruction naming

Example:

`movaps`

`mov` = move

`u` = unaligned

`p` = packaged

`s` = single precision

avx/sse/avx512

different architectures provide different instruction sets (each new version introduced more)

so older architectures do not have all instructions

mention double precision instruction (when added)

# What makes my code eligible for vectorization?

- calculations over arrays
- code must be in the innermost loop
- no if statements
- no uninline function calls

# How can I use vectorization?

The compiler does that for us if we tell him to.

Example for gcc:

- gcc standard optimizations does not vectorize
- -O3 enables auto vectorization
- -O3 does it by using the -ftree-vectorize flag
- -fopt-info-vec enables vectorization report



```
void test(float * vec1, float * vec2, float * res) {  
    for (unsigned long i = 0; i < vector_size; i++) {  
        res[i] += vec2[i] * vec1[i];  
        res[i] /= vec2[i];  
        res[i] -= vec1[i];  
    }  
}
```

- TODO add restrict explanation(compiler checks for overlapping arrays -> more assem code) restrict can be dangerous since it trusts the programmer to be right. use with caution. While explaining show assembler code

```
void test(float *__restrict vec1, float *__restrict vec2, int n) {  
    for (unsigned long i = 0; i < vector_size; i++) {  
        res[i] += vec2[i] * vec1[i];  
        res[i] /= vec2[i];  
        res[i] -= vec1[i];  
    }  
}
```

TODO add explanation for alignment (compiler does not know automatically how long a type is and checks for it resulting in more assembler code)

```
typedef float_32 attribute((aligned(32)))  
void test(float_32 *__restrict vec1,  
          float_32 *__restrict vec2,  
          float_32 *__restrict res)  
{  
    for (unsigned long i = 0; i < vector_size; i++) {  
        res[i] += vec2[i] * vec1[i];  
        res[i] /= vec2[i];  
        res[i] -= vec1[i];  
    }  
}
```

# Zusammenfassung

- Zusammenfassung 1
  - Unterpunkt 1
  - Unterpunkt 2
- Zusammenfassung 2
  - Unterpunkt 1
  - Unterpunkt 2
- Quelle: [?]

# Literatur